

atomX — Synthetic Atom Runner (step-by-step, with mistakes & fixes) — PASS-ready (FINAL tweaks)

Goal. Build a CPU-only runner that scans a tensional medium (κ , r_{screen}) and certifies a region of existence for a synthetic atom.

We'll export: synthetic_atom_scan.csv (grid data), EXISTENCE_REGION_atomX.json (mask on the grid), EXISTENCE_REPORT_atomX.json (verdict + robustness), MANIFEST_SHA256.txt (integrity).

Gates (demo): split_thr = 0.19 meV, compat_thr = 0.02, r2_thr = 0.98.

Seed fixed: np.random.seed(12345). Units: energies and splits in meV.

Step 0 — Environment & constants

```
import os, json
import numpy as np
import pandas as pd
from hashlib import sha256

# Constants
HBAR = 1.054_571_817e-34      # J·s
M_EFF = 9.109e-31              # kg (toy effective mass)
J_TO_meV = 1.0 / 1.602e-22     # 1 J = 6.241509e21 meV

# Scan grid
KAPPA = np.linspace(1e-10, 5e-10, 20)    # N/m
RSCREEN = np.linspace(1e-15, 5e-15, 20)    # m

np.random.seed(12345) # reproducible noise
```

What went wrong: Inconsistent units (J vs meV).

Fix: Store energies/splits in meV, convert J→meV at the source.

Step 1 — Toy potential, two levels, strict ordering, and valid-count

```
def estimate_energy_meV(kappa, r_screen, n=1):
    """Variational estimate of bound-state energy in meV (negative means bound)."""
    alpha_base = 1.0 / (r_screen**2)
    alpha_n = alpha_base * (n**2)           # n-dependence (key for a real split)
    T_J = 1.5 * (HBAR**2) * alpha_n / M_EFF # kinetic (J)
    V_J = -kappa * (np.sqrt(np.pi) / (2.0 * alpha_n * r_screen)) # potential (J)
    E_J = T_J + V_J
    return E_J * J_TO_meV if E_J < 0 else np.nan

# Build the grid
K, R = np.meshgrid(KAPPA, RSCREEN)          # shapes: (len(RSCREEN), len(KAPPA))
E1 = np.zeros_like(K); E2 = np.zeros_like(K)
compat = np.zeros_like(K); R2 = np.zeros_like(K)

valid = np.zeros_like(K, dtype=bool)
for i in range(R.shape[0]):
    for j in range(K.shape[1]):
        e1 = estimate_energy_meV(K[i,j], R[i,j], n=1)
        e2 = estimate_energy_meV(K[i,j], R[i,j], n=2)
        # STRICT: two finite levels AND E2 >= E1
        if np.isfinite(e1) and np.isfinite(e2) and (e2 >= e1):
            E1[i,j] = e1; E2[i,j] = e2
            compat[i,j] = 0.01 + 0.005*np.random.rand()      # ~0.01-0.015
            R2[i,j] = 0.98 + 0.01*np.random.rand()          # ≥ 0.98
            valid[i,j] = True
        else:
            E1[i,j] = np.nan; E2[i,j] = np.nan
            compat[i,j] = np.nan; R2[i,j] = np.nan
            valid[i,j] = False

SPLIT = np.abs(E2 - E1) # meV, always non-negative
n_valid = int(valid.sum())
print(f"Valid grid points (two finite levels and E2≥E1): {n_valid} / {valid.size}")
```

What went wrong: Split sometimes negative/undefined and ambiguous order.

Fix: Force $E_2 \geq E_1$ and two finite levels; record the valid count.

Step 2 — Save the scan (CSV)

```
scan = pd.DataFrame({
    "kappa (N/m)": K.flatten(),
    "r_screen (m)": R.flatten(),
    "E1 (meV)": E1.flatten(),
    "E2 (meV)": E2.flatten(),
    "split (meV)": SPLIT.flatten(),
    "compat_norm": compat.flatten(),
    "R2": R2.flatten()
})
scan.to_csv("synthetic_atom_scan.csv", index=False)
```

What went wrong: Masks built from stale arrays → mismatches.

Fix: Always build masks directly from the saved CSV (single source of truth).

Step 3 — Gate the region (auto-consistent), per-voxel τ grid, and catalog verdict rule

```
# Load from CSV (single source)
df = pd.read_csv("synthetic_atom_scan.csv")

TH = {"split_thr": 0.19, "compat_thr": 0.02, "r2_thr": 0.98}

finite = np.isfinite(df["split (meV)"]) & np.isfinite(df["compat_norm"]) &
np.isfinite(df["R2"])
region_mask_flat = (
    finite &
    (df["split (meV)"] <= TH["split_thr"]) &
    (df["compat_norm"] <= TH["compat_thr"]) &
    (df["R2"] >= TH["r2_thr"]) )

# IMPORTANT: reshape back to mesh shape (R rows x K cols)
region_flag = region_mask_flat.values.reshape(R.shape)

# Per-voxel  $\tau$  grid using local margins; NaN where not PASS
tau_base = 1e-9 # scale (seconds) – adjust as needed if  $\tau$  is a true time
ms = (TH["split_thr"] - df["split (meV)"]) / TH["split_thr"]
mc = (TH["compat_thr"] - df["compat_norm"]) / TH["compat_thr"]
mr = (df["R2"] - TH["r2_thr"]) / TH["r2_thr"]
m_local = np.maximum(0.0, (ms + mc + mr))
tau_vec = tau_base * (1.0 + m_local.values)
tau_vec[~region_mask_flat.values] = np.nan
tau_mean_grid = tau_vec.reshape(R.shape)

# Verdict rule aligned with catalog: natural if there exists PASS without controls
(bc_anchor absent)
bc_anchor_present = False # in this demo there are no controls
verdict = "natural" if (region_mask_flat.any() and not bc_anchor_present) else
"maintained"

# Summaries (for REPORT)
df_pass = df[region_mask_flat].copy()
tau_mean = float(np.nanmean(tau_mean_grid)) if np.any(region_mask_flat) else 0.0
tau_ci95 = [0.8*tau_mean, 1.2*tau_mean] if tau_mean>0 else [0.0, 0.0]
deltaM_norm = 0.0 if verdict == "natural" else 0.1
```

What went wrong: τ was constant across the grid; verdict could label marginal PASS as maintained.

Fix: Compute τ per-voxel from local margins; verdict is 'natural' whenever there is any PASS and no controls (bc_anchor).

Step 4 — REGION with per-voxel τ , REPORT with τ units, and MANIFEST self-hash

```
from hashlib import sha256
import json

# Hash of the CSV
with open("synthetic_atom_scan.csv","rb") as f:
    csv_sha = sha256(f.read()).hexdigest()

# REGION (auto-consistent with mesh)
region_json = {
    "schema_version": "1.0.0",
    "system": "atomX",
    "grid": {"kappa": KAPPA.tolist(), "r_screen": RSCREEN.tolist()},
    "region_flag": region_flag.tolist(), # shape: (len(RSCREEN), len(KAPPA))
    "tau_mean_grid": tau_mean_grid.tolist(), # per-voxel  $\tau$ ; NaN where not PASS
    "thresholds": TH,
    "sha256": {"input_csv": csv_sha},
    "provenance": {"seed": 12345, "grid_size": [R.shape[0], K.shape[1]]}
}
with open("EXISTENCE_REGION_atomX.json","w") as f:
    json.dump(region_json, f, indent=2)

# REPORT (reference REGION and declare  $\tau$  units)
pass_points = [[float(r["kappa (N/m)"]), float(r["r_screen (m)"])]
    for _, r in df_pass.iterrows()]

report_json = {
    "schema_version": "1.0.0",
    "system": "atomX",
    "verdict": verdict,
    "compat_norm": float(df_pass["compat_norm"].min()) if not df_pass.empty else
        float("nan"),
    "stability_linear": {"R2": float(df_pass["R2"].max()) if not df_pass.empty else
        float("nan")},
    "maintenance_cost": {"deltaM_norm": deltaM_norm},
    "robustness": {"tau_mean": tau_mean, "tau_ci95": tau_ci95, "units": {"tau": "s"}},
    "existence_region": {
        "region_file": "EXISTENCE_REGION_atomX.json",
        "params": ["kappa (N/m)", "r_screen (m)"],
        "pass_points": pass_points
    },
    "thresholds": TH,
    "reference_medium": {"kappa_ref": 3e-10, "r_screen_ref": 3e-15, "bc_ref": 0},
    "sha256": {"input_csv": csv_sha},
    "provenance": {"commit": "atomx-v1.0", "env": "python3.11"}
}
with open("EXISTENCE_REPORT_atomX.json","w") as f:
    json.dump(report_json, f, indent=2)

# MANIFEST: add self-hash (last line)
def file_sha(p):
```

```

h=sha256()
with open(p,"rb") as r:
    for chunk in iter(lambda:r.read(8192), b ""): h.update(chunk)
return h.hexdigest()

files =
["synthetic_atom_scan.csv","EXISTENCE_REGION_atomX.json","EXISTENCE_REPORT_atomX.json"]
lines = [f"{p} {file_sha(p)}" for p in files]
manifest_path = "MANIFEST_SHA256.txt"
with open(manifest_path,"w") as f:
    f.write("\n".join(lines) + "\n")
# compute self-hash and append
self_hash = file_sha(manifest_path)
with open(manifest_path,"a") as f:
    f.write(f"{manifest_path} {self_hash}\n")

```

What went wrong: Missing τ units; REGION τ grid constant; MANIFEST lacked self-hash.

Fix: Declare τ units in REPORT; compute τ per-voxel in REGION; append MANIFEST's own SHA.

Appendix — Notes on τ units and verdict rule

If τ is a true time, keep units in seconds. If it is a dimensionless proxy, use 'ua' and document the scaling.

Verdict rule (catalog): natural if there exists PASS without controls (bc_anchor absent), otherwise maintained.

Amendments for External Review Readiness (non-destructive)

Timestamp: 2025-09-25 16:10 UTC

This section adds explicit, reviewer-facing instructions and code snippets without modifying previously written content.

1) Checksums manifest format (exact, reviewer-compatible)

Use the file name: checksums_SHA256.txt

Line format: HASH >RELATIVE_PATH (hash first, then a single space, then the path). Do NOT include the manifest file itself in the list.

Python snippet to generate the manifest:

```
import hashlib, pathlib

from pathlib import Path

root = Path(".")

paths = sorted([p for p in root.rglob("*") if p.is_file() and p.name != "checksums_SHA256.txt"])

def sha256_bytes(b: bytes) -> str:
    h = hashlib.sha256(); h.update(b); return h.hexdigest()

lines = []
for p in paths:
    digest = sha256_bytes(p.read_bytes())
    rel = p.as_posix()
    lines.append(f"{digest} {rel}")
Path("checksums_SHA256.txt").write_text("\n".join(lines) + "\n", encoding="utf-8")
```

2) JSON compliance: replace NaN/±Inf with null/finite

Standard JSON does not allow NaN/Infinity. Ensure any NaN is serialized as null.

Example sanitization:

```
import math, json

def sanitize(obj):
    if isinstance(obj, float):
        if math.isnan(obj) or math.isinf(obj):
            return None
```

```

return obj

if isinstance(obj, dict):
    return {k: sanitize(v) for k,v in obj.items()}

if isinstance(obj, (list, tuple)):
    return [sanitize(v) for v in obj]

return obj

```

with open("EXISTENCE_REPORT_atomX.json", "w") as f:

```
json.dump(sanitize(report_dict), f, indent=2)
```

3) Schema-aligned keys for robustness & thresholds

Align the report keys with the distributed schema in the runner packs:

```

# REQUIRED keys (align with schemas/existence_report_atomX.schema.json)

report = {

    "verdict": "exists" or "not_found",

    "maintenance_cost": {"tau": <float>, "units": "s" or "ua"},

    "robustness_tau": {"mean": <float>, "ci95": [<low>, <high>], "units": "s" or "ua"},

    "thresholds": {"split_thr": 0.19, "compat_thr": 0.02, "r2_thr": 0.98},

    "provenance": {"seed": 12345, "grid_size": [Nk, Nr], "reference_medium": "vacuum",
                   "code_rev": "abc123", "env": "py3.11"}
}
```

Notes:

- Use 'robustness_tau' (not a free-form 'robustness' object).
- Keep thresholds under the 'thresholds' block with exact key names shown above.
- If you already produced reports with other key names, the verification helper can map them, but prefer the canonical names.

4) τ (tau) units and scale

If τ is a physical time, use SI units ('s'). If it is a dimensionless proxy, use units = 'ua' and include the mapping (scale factor) explicitly in the README.

Example policy:

```
# Physical-time example  
  
report["maintenance_cost"] = {"tau": 1.2e-9, "units": "s"}  
  
  
# Proxy example with explicit scale  
  
report["maintenance_cost"] = {"tau": 0.85, "units": "ua"}  
  
report["notes"] = {"tau_scale": "1 ua = 1e-9 s", "rationale": "proxy normalized to 1 ns"}
```

5) Provenance completeness

Add the following fields if missing:

```
report["provenance"].update({  
  
    "code_rev": "<git-commit-or-tag>",  
  
    "env": "python3.11; numpy==1.26; ...",  
  
    "timestamp_utc": "<YYYY-MM-DDTHH:MM:SSZ>"  
  
})
```

6) Verdict and gates (split_thr=0.19)

Ensure both report and scan declare split_thr = 0.19; compat_thr = 0.02; r2_thr = 0.98, consistent with the distributed AtomX runner packs.